

Saksham Adhikari

737-315-1963 | adhsaksham27@gmail.com | linkedin.com/in/adhsaksham | github.com/Tar-ive

EDUCATION

Texas State University <i>B.S. in Computer Information Systems, AI Concentration, GPA: 3.83/4.0</i>	San Marcos, TX May 2027
• Coursework: Data Structures & Algorithms, Database Systems, Machine Learning, Data Mining in R • Achievements: 6x Hackathon Winner; Most Commercializable Hack – AWS Hackathon; Winner – AITX NVIDIA Hackathon; Gold Medal, Nepal Economics Olympiad; \$48k in scholarships; Google TPU Student Researcher Grant	

PROFESSIONAL EXPERIENCE

Machine Learning Engineer Intern <i>AskSLM</i>	May 2026 – Aug 2026 Austin, TX
• Quantized custom-trained LLaVA vision models for edge devices, monitoring weight clustering and achieving a 3x increase in image-to-text sequence recognition throughput • Enhanced internal Retrieval-Augmented Generation (RAG) pipeline by implementing Reciprocal Rank Fusion, boosting retrieval precision and achieving a 3% improvement in Mean Reciprocal Rank (MRR) • Engineered real-time video analytics pipelines via DeepStream SDK, optimizing frame-rate throughput and end-to-end inference latency using Harbor for systematic evaluation	
Software Engineer Intern <i>AskSLM</i>	Jan 2026 – May 2026 Austin, TX
• Fine-tuned a LLaMA 8B model with 3,000 curated QA pairs from internal documentation, integrating a RAG architecture for factual grounding and securely deploying via mTLS encryption • Migrated and optimized the internal vLLM serving stack to fit NVIDIA DGX Spark VRAM constraints, redesigning inference configuration to achieve a 2.4x throughput increase while operating within hardware memory limits • Integrated the RAG agent with Slack for daily utilization on NVIDIA DGX Spark devices, reducing developer documentation search time by 30%	
Research Assistant <i>Translational Health Research Center</i>	Jul 2025 – Nov 2025 San Marcos, TX
• Co-authored 2 published papers, building a data pipeline that reconciled U.S. Census datasets containing ~1.2M rows • Identified 7 underserved geographic clusters and modeled disease spread across 260,000 patients using multivariate and Graph Neural Network models; presented findings to non-technical partners	
Software Development Intern <i>Breast Cancer Resource Center</i>	May 2025 – Jul 2025 Austin, TX
• Built a React patient-facing chat agent backed by a RAG pipeline, reducing caregiver workload by 20% while improving response grounding and hallucination reduction • Improved CSV processing speeds from 1900ms to 1200ms using asynchronous coroutines, enabling faster billing calculations while adhering to internal compliance requirements	
Student Software Developer <i>Texas State Center for Analytics and Data Science</i>	Dec 2024 – May 2025 San Marcos, TX
• Built a hybrid RAG pipeline pairing FAISS dense vector search with cosine-similarity ranking for AI-driven grant proposal optimization, reducing faculty evaluation time by 80% • Managed the TXST grant pipeline on a Red Hat Enterprise Linux test server via a custom GitLab CI/CD pipeline with linting and dependency checks gating every merge; optimized CI/CD pipeline overhead, cutting build-and-deploy time by 21% to under 5 minutes	

PROJECTS

Giftmaxxing <i>Django, Swift, Terraform</i>	2026
• Won \$6k+ and the Most Commercializable Hack at the AWS v0 Hackathon, then iterated the application toward release • Engineered a shopping agent and trained a two-tower ranking model with real-time item retrieval	
TorchLeet <i>PyTorch, Deep Neural Networks, Tensor Math</i>	2026
• Core contributor to TorchLeet (LeetCode for PyTorch, 2.5k stars): 6 PRs merged across frontend fixes and technical answers	

LEADERSHIP

ACM AI @ TXST – Vice President	December 2024 – August 2025
• Grew the community of AI developers on campus to 100+ people and organized 10+ technical sessions on MCPs, CLI tools, and OpenClaw	

TECHNICAL SKILLS

Languages: Python, Go, Swift, R, Bash, JavaScript, SQL
AI/ML: PyTorch, TensorFlow, Scikit-learn, LLMs, RAG, Fine-tuning (LoRA, QLoRA), Quantization, Distillation, RRF, FAISS
Systems & Infra: Distributed Systems, Concurrency, GCP, AWS, Docker, Kubernetes, CI/CD, Linux, vLLM, DeepStream